

# BigLake

Built on years of investment in BigQuery, BigLake is a storage engine that allows organizations to unify data warehouses and lakes, and enable them to perform uniform fine-grained access control, and accelerate query performance across multi-cloud storage and open formats.

- Store a single copy of data with uniform features across data warehouses & lakes.
- Fine-grained access control and multi-cloud governance over distributed data.
- Seamless integration with open source analytics tools and open data formats.

## BENEFITS

### **Freedom of choice**

Unlock analytics on distributed data regardless where and how it's stored, while choosing the best analytics tools, open source or cloud native over a single copy of data.

### **Secure and performant data lakes**

Fine-grained access control across open source engines like Apache Spark, Presto and Trino, and open formats such as Parquet. Performant queries over data lakes powered by BigQuery.

### **Unified governance & management at scale**

Integrates with [Dataplex](#) to provide management at scale, including logical data organization, centralized policy & metadata management, quality and lifecycle management for consistency across distributed data.

## KEY FEATURES

### **Key features**

#### **Fine grained security controls**

BigLake eliminates the need to grant file level access to end users. Apply table, row, column level security policies on object store tables similar to existing BigQuery tables.

**Multi-compute analytics**

Maintain a single copy of data and make it uniformly accessible across Google Cloud and open-source engines, including [BigQuery](#), [Vertex AI](#), [Dataflow](#), Spark, Presto, Trino, and Hive using BigLake connectors. Centrally manage security policies in one place, and have it consistently enforced across the query engines by the API interface built into the connectors.

**Multi-cloud governance**

Discover all BigLake tables, including those defined over Amazon S3, Azure data lake Gen 2 in [Data Catalog](#). Configure fine grained access control and have it enforced across clouds when querying with [BigQuery Omni](#).

**Performance acceleration**

Achieve industry leading performance over data lake tables on Google Cloud, AWS and Azure, powered by proven BigQuery infrastructure.

**Built on open formats**

Gain access to the most popular open data formats including Parquet, Avro, ORC, CSV, JSON. The API serves multiple compute engines through Apache Arrow.