





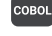



























-  Secrets
-  ABAP
-  Apex
-  C
-  C++
-  CloudFormation
-  COBOL
-  C#
-  CSS
-  Flex
-  Go
-  HTML
-  Java
-  JavaScript
-  Kotlin
-  **Kubernetes**
-  Objective C
-  PHP
-  PL/I
-  PL/SQL
-  Python
-  RPG
-  Ruby
-  Scala
-  Swift
-  Terraform
-  Text
-  TypeScript
-  T-SQL
-  VB.NET
-  VB6
-  XML



Kubernetes static code analysis


Unique rules to find Security Hotspots in your KUBERNETES code


- All rules 7
-  Security Hotspot 6
-  Code Smell 1


Tags ▾


Search by name... 🔍


- Mounting sensitive file system paths is security-sensitive


 Security Hotspot
- Using host operating system namespaces is security-sensitive


 Security Hotspot
- Allowing process privilege escalations is security-sensitive

 Security Hotspot
- Exposing Docker sockets is security-sensitive

 Security Hotspot
- Running containers in privileged mode is security-sensitive

 Security Hotspot
- Setting capabilities is security-sensitive

 Security Hotspot
- Kubernetes parsing failure

 Code Smell

Allowing process privilege escalations is security-sensitive

Analyze your code

-  Security Hotspot
-  Major ?
-  cwe

Allowing process privilege escalations exposes the Pod to attacks that exploit setuid binaries.

This field directly controls whether the `no_new_privs` flag is set in the container process.

When this flag is enabled, binaries configured with `setuid` or `setgid` bits cannot change their runtime uid or gid: Potential attackers must rely on other privilege escalation techniques to successfully operate as root on the Pod.

Depending on how resilient the Kubernetes cluster and Pods are, attackers can extend their attack to the cluster by compromising the nodes from which the cluster started the Pod.

The `allowPrivilegeEscalation` field should not be set to `true` unless the Pod's risks related to `setuid` or `setgid` bits have been mitigated.

Ask Yourself Whether

- This Pod is accessible to people who are not administrators of the Kubernetes cluster.
- This Pod contains binaries with `setuid` or `setgid` capabilities.

There is a risk if you answered yes to all of these questions.

Recommended Secure Coding Practices

Disable privilege escalation.

Sensitive Code Example

```
apiVersion: v1
kind: Pod
metadata:
  name: example
spec:
  containers:
    - name: web
      image: nginx
      ports:
        - name: web
          containerPort: 80
          protocol: TCP
      securityContext:
        allowPrivilegeEscalation: true # Sensitive
```

Compliant Solution

```
apiVersion: v1
kind: Pod
metadata:
  name: example
spec:
  containers:
    - name: web
      image: nginx
      ports:
```

```
- name: web
  containerPort: 80
  protocol: TCP
securityContext:
  allowPrivilegeEscalation: false
```

See

- [MITRE, CWE-284](#) - Improper Access Control
- [Linux Kernel Archives, no_new_privs](#) - Official docs

Available In:

