

Test model updates with reproducible training

cover, presence of a rainbow, and several other features.

Perhaps you want to continue improving your unicorn model. For example, suppose you do some additional feature engineering on a certain feature and then retrain the model, hoping to get better (or at least the same) results. Unfortunately, it is sometimes difficult to reproduce model training. To improve reproducibility, follow these recommendations:

- Deterministically seed the random number generator. For details, see randomization in data generation
- Initialize model components in a fixed order to ensure the components get the same random number from the random number generator on every run. ML libraries typically handle this requirement automatically.
- Take the average of several runs of the model.

∓ Filter

overniting (105 min)

Advanced ML models

▶ Neural networks (75 min)

■ Introduction (2 min)

► Embeddings (45 min)

Real-world ML

min)

min)

(10 min)

→ What's next

► Fairness (110 min)

min)

• Use version control, even for preliminary iterations, so that you can pinpoint code and parameters when investigating your model or pipeline.

Even after following these guidelines, other sources of nondeterminism might still exist.

Test calls to machine learning API

How do you test updates to API calls? You could retrain your model, but that's time intensive. Instead, write a unit test to generate random input data and run a single step of gradient descent. If this step completes without errors, then any updates to the API probably haven't ruined your model.

Write integration tests for pipeline components

In an ML pipeline, changes in one component can cause errors in other components. Check that components work together by writing an integration test that runs the entire pipeline end-to-end.

Besides running integration tests continuously, you should run integration tests when pushing new models and new software versions. The slowness of running the entire pipeline makes continuous integration testing harder. To run integration tests faster, train on a subset of the data or with a simpler model. The details depend on your model and data. To get continuous coverage, you'd adjust your faster tests so that they run with every new version of model or software. Meanwhile, your slow tests would run continuously in the background.

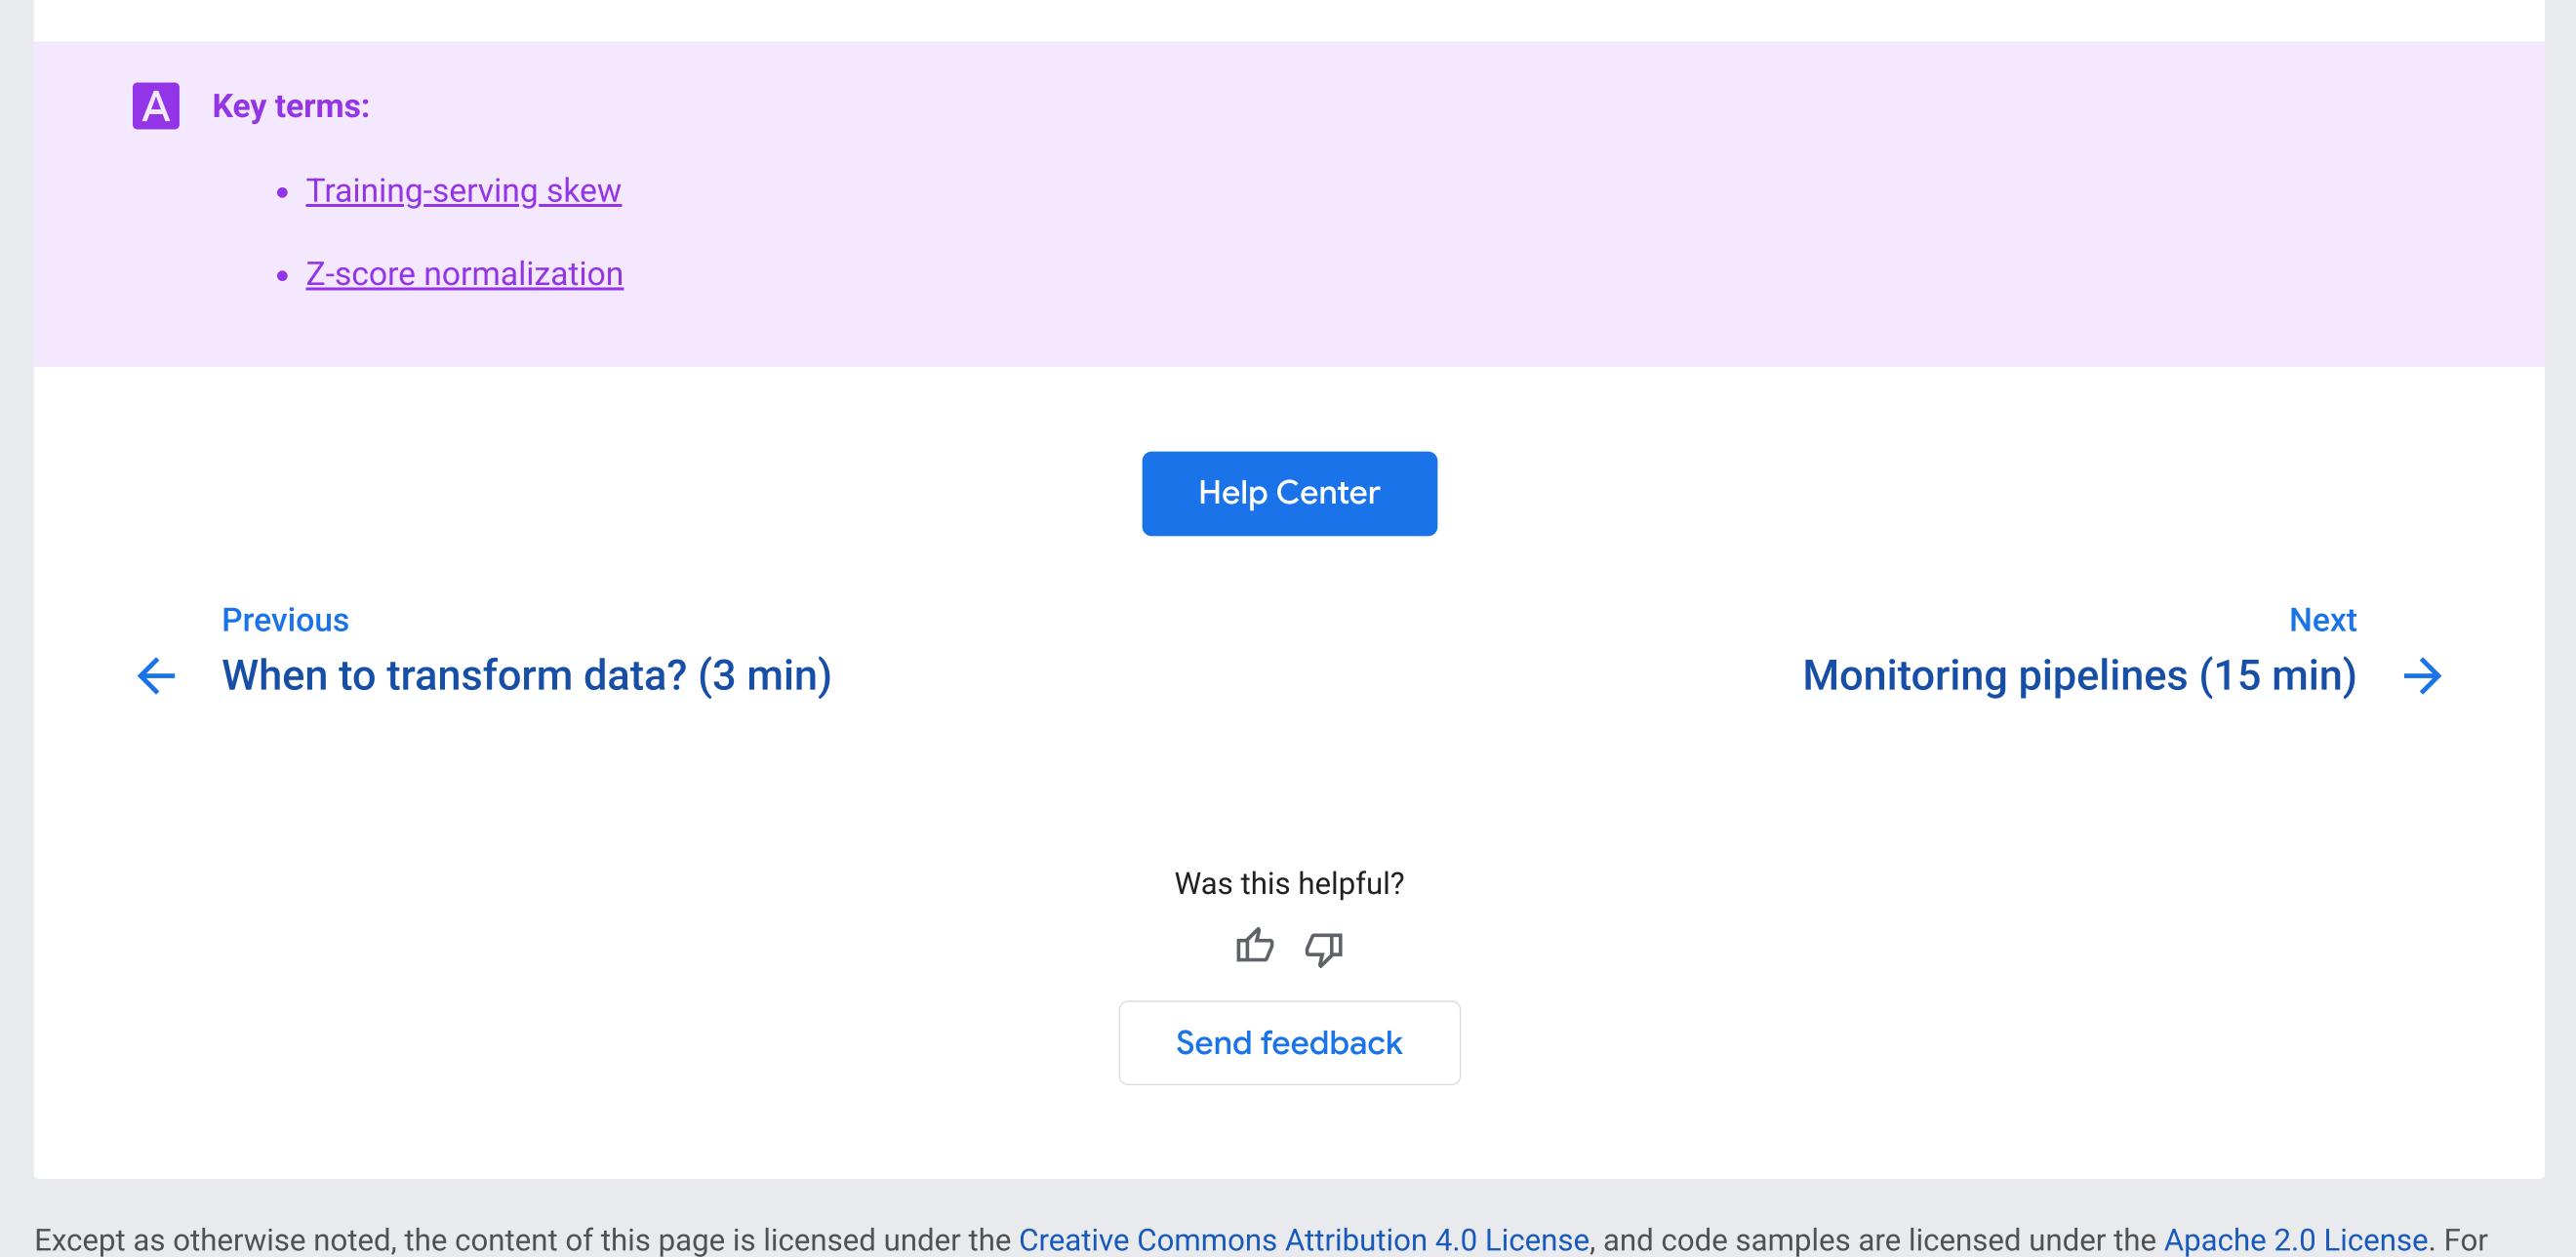
Validate model quality before serving

Before pushing a new model version to production, test for the following two types of quality degradations:

- Sudden degradation. A bug in the new version could cause significantly lower quality. Validate new versions by checking their quality against the previous version.
- Slow degradation. Your test for sudden degradation might not detect a slow degradation in model quality over multiple versions. Instead, ensure your model's predictions on a validation dataset meet a fixed threshold. If your validation dataset deviates from live data, then update your validation dataset and ensure your model still meets the same quality threshold.

Validate model-infrastructure compatibility before serving

If your model is updated faster than your server, then your model could have different software dependencies from your server, potentially causing incompatibilities. Ensure that the operations used by the model are present in the server by staging the model in a sandboxed version of the server.



details, see the Google Developers Site Policies. Java is a registered trademark of Oracle and/or its affiliates.

Last updated 2024-10-09 UTC.

