

ML Concepts

Home Crash Course

Filter

Advanced ML models

- Neural networks (75 min)
- Embeddings (45 min)
- Large language models (LLMs) (45 min)

Real-world ML

- Production ML systems (70 min)
- Automated machine learning (30 min)
- Fairness (110 min)
  - Introduction (5 min)
  - Types of bias (5 min)
  - Identifying bias (10 min)
  - Mitigating bias (5 min)
  - Evaluating for bias (5 min)
  - Demographic parity (10 min)
  - Equality of opportunity (10 min)
  - Counterfactual fairness (10 min)



Home > Products > Machine Learning > ML Concepts > Crash Course

Was this helpful?

# Fairness



Send feedback

Estimated module length: 110 minutes

Evaluating a machine learning model (ML) responsibly requires doing more than just calculating overall loss metrics. Before putting a model into production, it's critical to audit training data and evaluate predictions for [bias](#).

This module looks at different types of human biases that can manifest in training data. It then provides strategies to identify and mitigate them, and then evaluate model performance with fairness in mind.

## Learning objectives

- Become aware of common human biases that can inadvertently be reproduced by ML algorithms.
- Proactively explore data to identify sources of bias before training a model.
- Evaluate model predictions for bias.

## Prerequisites:

This module assumes you are familiar with the concepts covered in the following modules:

- [Introduction to Machine Learning](#)
- [Linear regression](#)
- [Logistic regression](#)
- [Classification](#)
- [Working with numerical data](#)
- [Working with categorical data](#)
- [Datasets, generalization, and overfitting](#)

Google for Developers  
2.51M subscribers

Crash Course: Fairness

Share

FAIRNESS

Watch on

## Key terms:

- [Bias \(ethics/fairness\)](#)
- [Model](#)

Help Center

[Previous](#)  
← **Getting started (10 min)**

**Types of bias (5 min)** → [Next](#)

Was this helpful?



Send feedback

Except as otherwise noted, the content of this page is licensed under the [Creative Commons Attribution 4.0 License](#), and code samples are licensed under the [Apache 2.0 License](#). For details, see the [Google Developers Site Policies](#). Java is a registered trademark of Oracle and/or its affiliates.

Last updated 2025-01-03 UTC.

Connect

Blog

Instagram

LinkedIn

X (Twitter)

YouTube

Programs

Google Developer Groups

Google Developer Experts

Accelerators

Women Techmakers

Google Cloud & NVIDIA

Developer consoles

Google API Console

Google Cloud Platform Console

Google Play Console

Firebase Console

Actions on Google Console

Cast SDK Developer Console

Chrome Web Store Dashboard

Google Home Developer Console

Google for Developers

Android

Chrome

Firebase

Google Cloud Platform

Google AI

All products

Terms | Privacy

Sign up for the Google for Developers newsletter

Subscribe

English