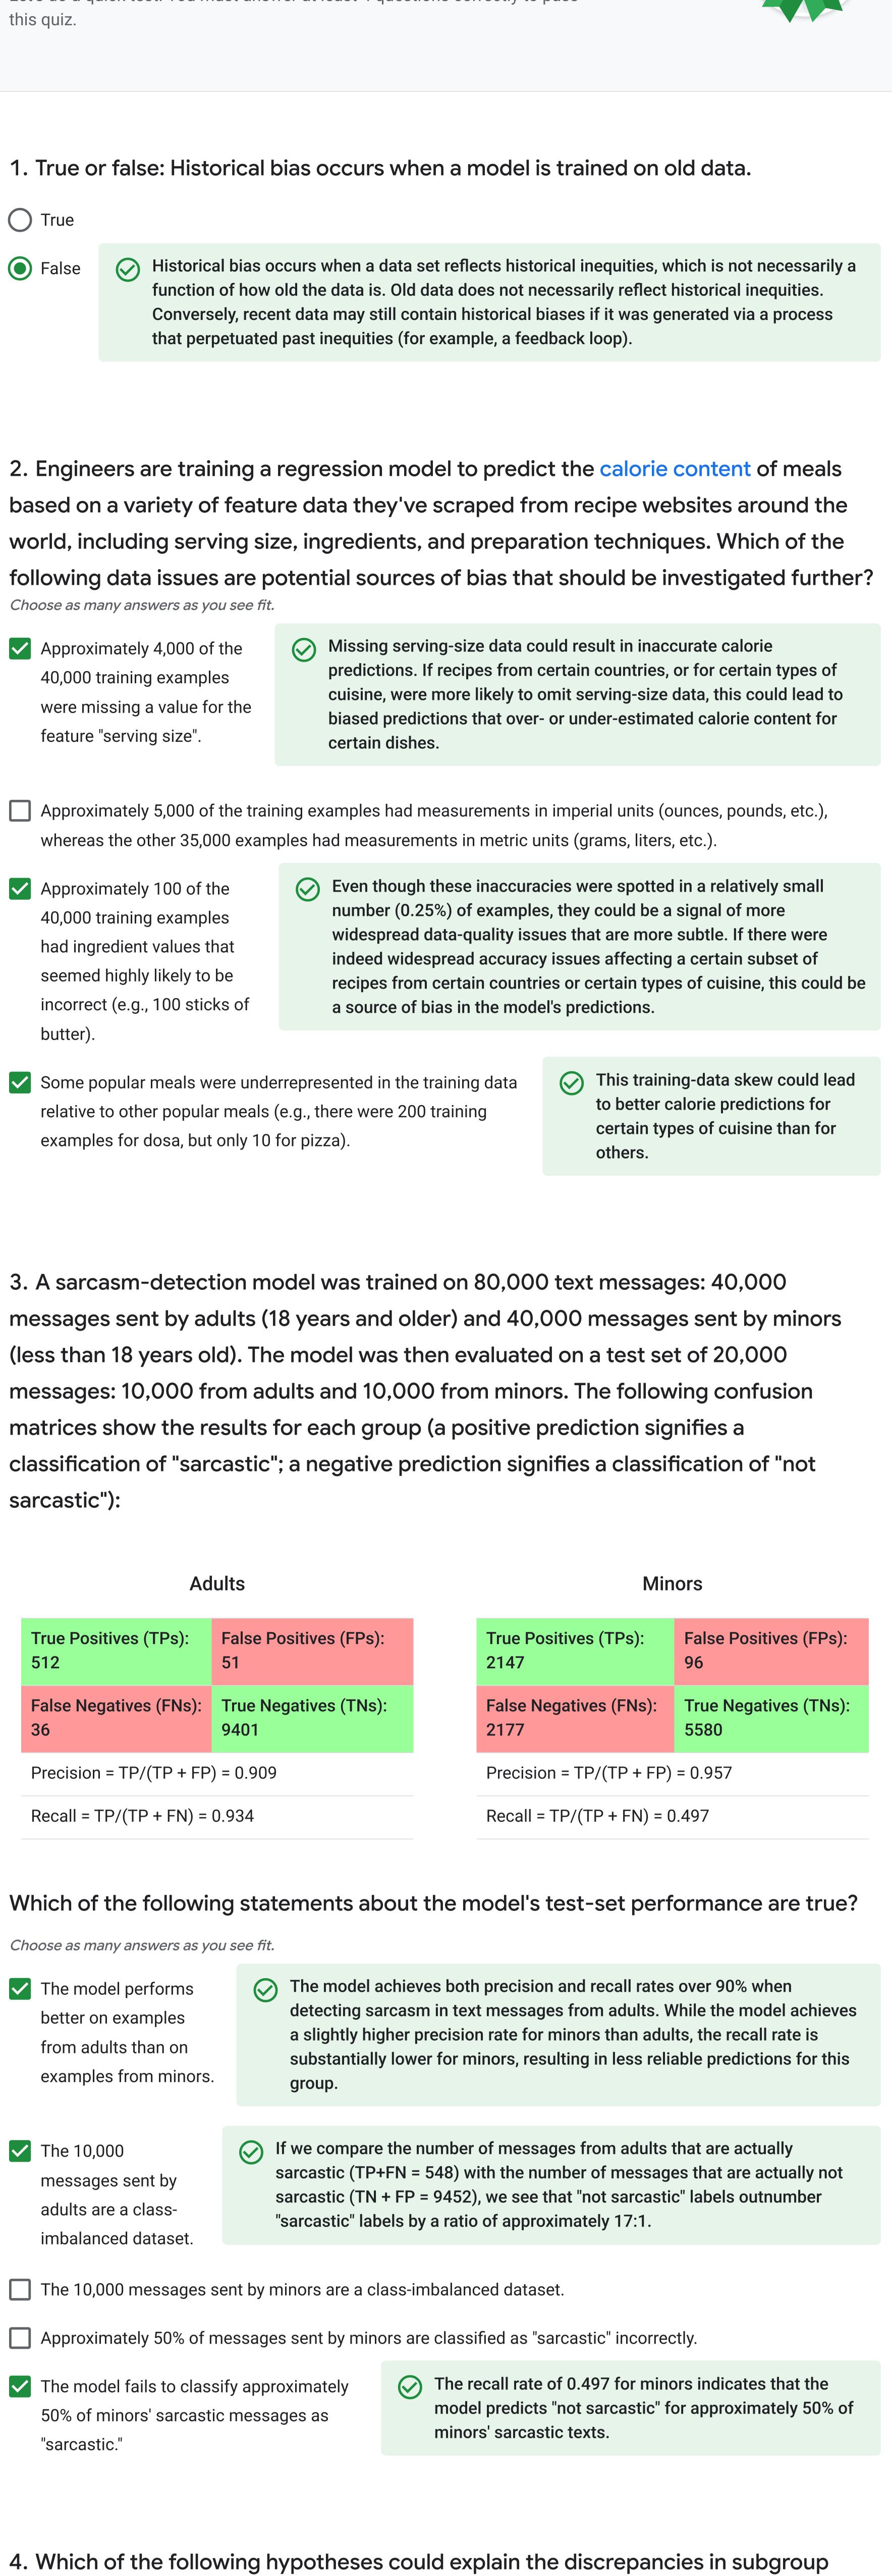
Let's do a quick test! You must answer at least 4 questions correctly to pass this quiz.





Choose as many answers as you see fit. The model errs too much on the side of predicting "sarcastic." As a result, it makes more errors when classifying minors' text messages, because there are more sarcastic messages from minors in the test set.

performance on the test set for the sarcasm-detection model above?

The model was evaluated on more negative (not-sarcastic) examples from minors than from adults, resulting in more errors for minors.

and thus less likely to be flagged by the model. There are far fewer actual sarcastic messages from adults than from

messages was more subtle,

Sarcasm in minors' text

minors' messages than in adults' messages, that could explain the discrepancy in recall rates between the two groups. The subset of text messages from adults is classimbalanced; Approximately 1 in 20 messages is actually

The sarcasm-detection model classified nearly 50% of minors'

sarcastic messages as "not sarcastic." If sarcasm was more subtle in

minors. If the model were evaluated on a more class-balanced set of adult messages, its recall might drop for that subgroup.

the model's inability to generalize, and that recall for text messages from adults might drop if the model were evaluated on addditional batches of adults' sarcastic messages.

sarcastic. It's possible that this class imbalance is masking

model has already been released into production. Which of the following stopgap strategies will help mitigate errors in the model's predictions? Restrict the model's usage to text messages sent by minors.

Adjust the model output so that it returns "sarcastic" for all text messages sent by minors, regardless of what

inconsistencies in sarcasm-detection accuracy across age demographics, but the

5. Engineers are working on retraining the sarcasm model above to address

the model originally predicted. The precision rate for text messages sent by minors is high, which means When the model predicts

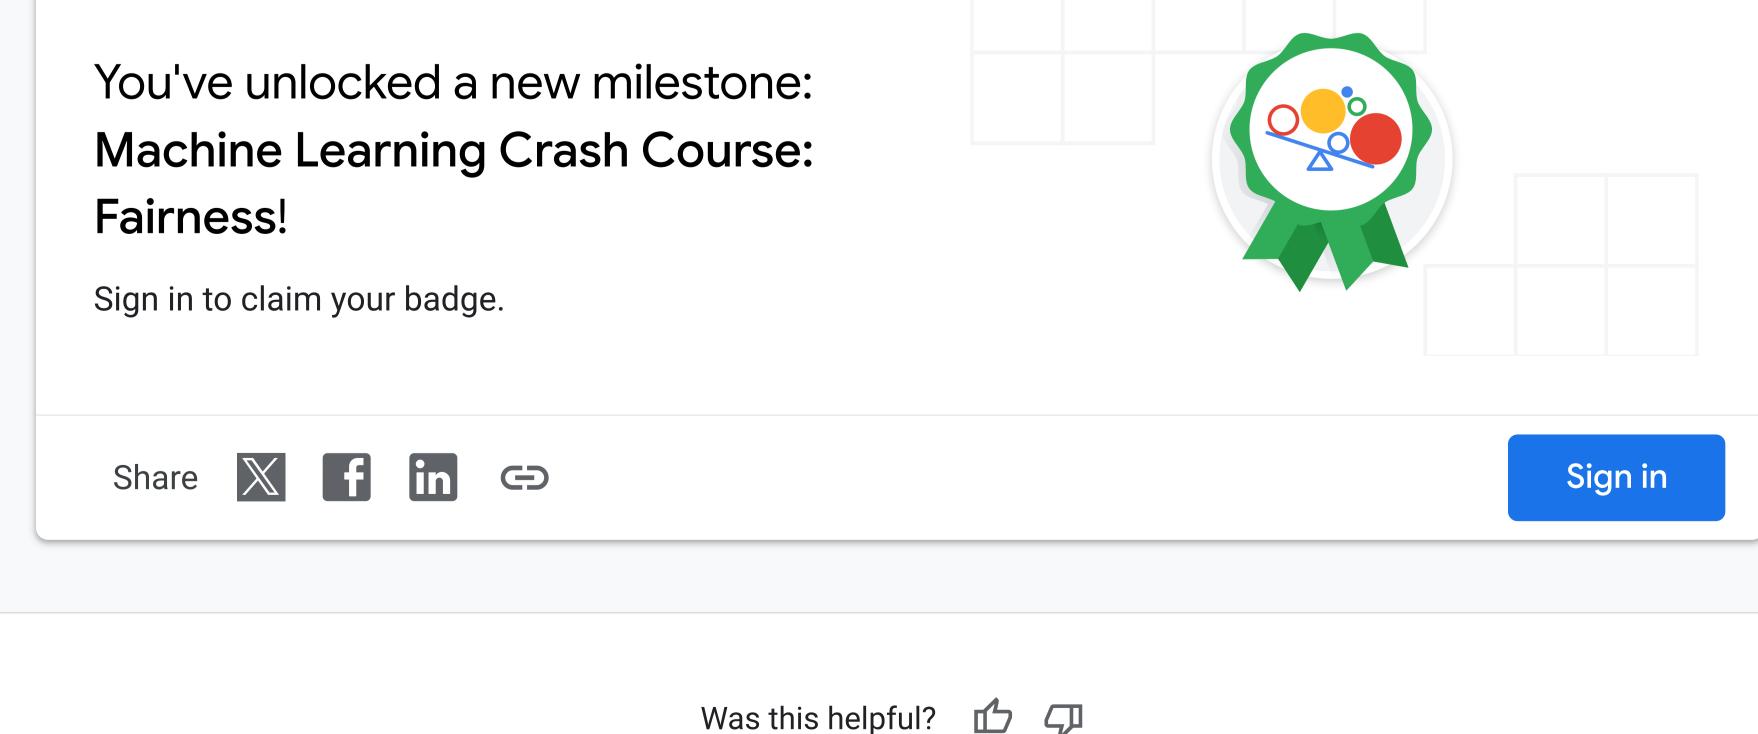
messages sent by minors, adjust the output so the model returns a value of "unsure" instead.

"not sarcastic" for text

correct. The problem is that recall is very low for minors; the model fails to identify sarcasm in approximately 50% of examples. Given that the model's negative predictions for minors are no better than random guesses, we can avoid these errors by not providing a prediction in these cases.

that when the model predicts "sarcastic" for this group, it is nearly always

Results You scored 5 out of 5. Congratulations! You have passed this quiz.



Google Al

All products

Connect Programs Blog Google Developer Groups Google Developer Experts Instagram LinkedIn Accelerators X (Twitter) Women Techmakers YouTube Google Cloud & NVIDIA

Chrome

Android

Developer consoles Google API Console Google Cloud Platform Console Google Play Console Firebase Console Actions on Google Console Cast SDK Developer Console Chrome Web Store Dashboard Google Home Developer Console

Firebase

Google Cloud Platform

Google for Developers